



REINFORCEMENT LEARNING

CP8319/CPS824

Lecture 10

Instructor: Nariman Farsad

Today's Agenda

- 1. Review of Previous Lectures**
2. Model-Free Control (Monte Carlo)
3. Model-Free Control (Temporal Difference)
4. Model-Free Control (Q-Learning)


Policy Iteration: Known Model

Set $i = 0$

Initialize $\pi_0(s)$ randomly for all states s

While $i \neq 0$ or $\|\pi_i - \pi_{i-1}\|_1 > 0$ (L1-norm, measures if the policy changed for any state):

- $v^{\pi_i} \leftarrow$ MDP value function **policy evaluation** of π_i
- $\pi_{i+1} \leftarrow$ **Policy improvement** on v^{π_i}
- $i = i + 1$



Policy evaluation Estimate v_π

Iterative policy evaluation

Policy improvement Generate $\pi' \geq \pi$

Greedy policy improvement

Model Free Policy Iteration

Set $i = 0$

Initialize $\pi_0(s)$ randomly for all states s

While $i \neq 0$ or $\|\pi_i - \pi_{i-1}\|_1 > 0$ (L1-norm, measures if the policy changed for any state):


- $Q^{\pi_i} \leftarrow$ MDP value function **MC policy Q evaluation** of π_i
- $\pi_{i+1} \leftarrow$ **ϵ -greedy Policy improvement** on Q^{π_i}
- $i = i + 1$

greedy(Q)

$$\pi_{i+1}(a|s) = \begin{cases} 1, & \text{if } a = \operatorname{argmax}_{a' \in \mathcal{A}} Q^{\pi_i}(s, a') \\ 0, & \text{otherwise} \end{cases}$$

ϵ -greedy(Q)

$$\pi_{i+1}(a|s) = \begin{cases} \epsilon/m + 1 - \epsilon, & \text{if } a = \operatorname{argmax}_{a' \in \mathcal{A}} Q^{\pi_i}(s, a') \\ \epsilon/m, & \text{otherwise} \end{cases}$$



Policy evaluation Monte-Carlo policy evaluation, $Q = q_\pi$

Policy improvement ϵ -greedy policy improvement

ϵ -Greedy and Greedy Example

- Let's say in a state s_1 we can take 3 actions a_1, a_2, a_3 .
- Assume that $a_1 = \operatorname{argmax}_{a' \in \mathcal{A}} Q^{\pi_i}(s, a')$ and $\epsilon = 0.5$.
- What are the ϵ -greedy and greedy policies $\pi(a|s_1)$?

greedy(Q)

$$\pi_{i+1}(a|s) = \begin{cases} 1, & \text{if } a = \operatorname{argmax}_{a' \in \mathcal{A}} Q^{\pi_i}(s, a') \\ 0, & \text{otherwise} \end{cases}$$

ϵ -greedy(Q)

$$\pi_{i+1}(a|s) = \begin{cases} \epsilon/m + 1 - \epsilon, & \text{if } a = \operatorname{argmax}_{a' \in \mathcal{A}} Q^{\pi_i}(s, a') \\ \epsilon/m, & \text{otherwise} \end{cases}$$

Today's Agenda

1. Review of Previous Lectures
- 2. Model-Free Control (Monte Carlo)**
3. Model-Free Control (Temporal Difference)
4. Model-Free Control (Q-Learning)

Model Free Policy Iteration

Set $i = 0$

Initialize $\pi_0(s)$ randomly for all states s


While $i \neq 0$ or $\|\pi_i - \pi_{i-1}\|_1 > 0$ (L1-norm, measures if the policy changed for any state):

- $Q^{\pi_i} \leftarrow$ MDP value function **MC policy Q evaluation** of π_i
- $\pi_{i+1} \leftarrow$ **ϵ -greedy Policy improvement** on Q^{π_i}
- $i = i + 1$

ϵ -greedy(Q)

$$\pi_{i+1}(a|s) = \begin{cases} \epsilon/m + 1 - \epsilon, & \text{if } a = \operatorname{argmax}_{a' \in \mathcal{A}} Q^{\pi_i}(s, a') \\ \epsilon/m, & \text{otherwise} \end{cases}$$

Does ϵ -greedy step provably improve policy?



Policy evaluation Monte-Carlo policy evaluation, $Q = q_\pi$

Policy improvement ϵ -greedy policy improvement

Monotonic ϵ -Greedy Policy Improvement

Theorem

For any ϵ -greedy policy π_i , the ϵ -greedy policy w.r.t. Q^{π_i} , π_{i+1} is a monotonic improvement $V^{\pi_{i+1}} \geq V^{\pi_i}$

$$\begin{aligned} Q^{\pi_i}(s, \pi_{i+1}(s)) &= \sum_{a \in A} \pi_{i+1}(a|s) Q^{\pi_i}(s, a) \\ &= (\epsilon/|A|) \left[\sum_{a \in A} Q^{\pi_i}(s, a) \right] + (1 - \epsilon) \max_a Q^{\pi_i}(s, a) \\ &= (\epsilon/|A|) \left[\sum_{a \in A} Q^{\pi_i}(s, a) \right] + (1 - \epsilon) \max_a Q^{\pi_i}(s, a) \frac{1 - \epsilon}{1 - \epsilon} \\ &= (\epsilon/|A|) \left[\sum_{a \in A} Q^{\pi_i}(s, a) \right] + (1 - \epsilon) \max_a Q^{\pi_i}(s, a) \sum_{a \in A} \frac{\pi_i(a|s) - \frac{\epsilon}{|A|}}{1 - \epsilon} \\ &\geq \frac{\epsilon}{|A|} \left[\sum_{a \in A} Q^{\pi_i}(s, a) \right] + (1 - \epsilon) \sum_{a \in A} \frac{\pi_i(a|s) - \frac{\epsilon}{|A|}}{1 - \epsilon} Q^{\pi_i}(s, a) \\ &= \sum_{a \in A} \pi_i(a|s) Q^{\pi_i}(s, a) = V^{\pi_i}(s) \end{aligned}$$

Each step of policy improvement improves policy or keeps it the same

Model Free Policy Iteration

Set $i = 0$

Initialize $\pi_0(s)$ randomly for all states s


While $i \neq 0$ or $\|\pi_i - \pi_{i-1}\|_1 > 0$ (L1-norm, measures if the policy changed for any state):

- $Q^{\pi_i} \leftarrow$ MDP value function **MC policy Q evaluation** of π_i
- $\pi_{i+1} \leftarrow$ **ϵ -greedy Policy improvement** on Q^{π_i}
- $i = i + 1$

ϵ -greedy(Q)

$$\pi_{i+1}(a|s) = \begin{cases} \epsilon/m + 1 - \epsilon, & \text{if } a = \operatorname{argmax}_{a' \in \mathcal{A}} Q^{\pi_i}(s, a') \\ \epsilon/m, & \text{otherwise} \end{cases}$$

Although ϵ -greedy step provably improve policy, does it converge to the optimal policy?



Policy evaluation Monte-Carlo policy evaluation, $Q = q_\pi$

Policy improvement ϵ -greedy policy improvement

Greedy in the Limit of Infinite Exploration (GLIE)

Definition of GLIE

- All state-action pairs are visited an infinite number of times

$$\lim_{i \rightarrow \infty} N_i(s, a) \rightarrow \infty$$

- Behavior policy (policy used to act in the world) converges to greedy policy

$$\lim_{i \rightarrow \infty} \pi(a|s) \rightarrow \arg \max_a Q(s, a) \text{ with probability 1}$$

- A simple GLIE strategy is ϵ -greedy where ϵ is reduced to 0 with the following rate: $\epsilon_i = 1/i$

GLIE Monte Carlo Control

Theorem


GLIE Monte-Carlo control converges to the optimal state-action value function $Q(s, a) \rightarrow Q^*(s, a)$

Monte Carlo Online Control/On Policy Improvement

```
1: Initialize  $Q(s, a) = 0, N(s, a) = 0 \forall (s, a)$ , Set  $\epsilon = 1, k = 1$ 
2:  $\pi_k = \epsilon$ -greedy( $Q$ ) // Create initial  $\epsilon$ -greedy policy
3: loop
4:   Sample  $k$ -th episode  $(s_{k,1}, a_{k,1}, r_{k,1}, s_{k,2}, \dots, s_{k,T})$  given  $\pi_k$ 
4:    $G_{k,t} = r_{k,t} + \gamma r_{k,t+1} + \gamma^2 r_{k,t+2} + \dots + \gamma^{T-t} r_{k,T}$ 
5:   for  $t = 1, \dots, T$  do
6:     if First visit to  $(s, a)$  in episode  $k$  then
7:        $N(s, a) = N(s, a) + 1$ 
8:        $Q(s_t, a_t) = Q(s_t, a_t) + \frac{1}{N(s,a)} (G_{k,t} - Q(s_t, a_t))$ 
9:     end if
10:  end for
11:   $k = k + 1, \epsilon = 1/k$ 
12:   $\pi_k = \epsilon$ -greedy( $Q$ ) // Policy improvement
13: end loop
```

MC for On Policy Control Example

- Robot with two actions
 - $R(-, a_1) = [1 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 +10]$ and $R(-, a_2) = [0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 +5]$
- $\pi(s) = a_1 \ \forall s, \gamma = 1, \epsilon = 0.5$. Any action from s_1 and s_7 terminates episode
- Sample episode = $(s_3, a_1, 0, s_2, a_2, 0, s_3, a_1, 0, s_2, a_2, 0, s_1, a_1, 1 \text{ terminal})$
- First visit MC estimate of Q of each (s, a) pair?
 - $Q^\pi(-, a_1) = [1 \ 0 \ 1 \ 0 \ 0 \ 0 \ 0], \quad Q^\pi(-, a_2) = [0 \ 1 \ 0 \ 0 \ 0 \ 0 \ 0]$
- What is $\pi(s) = \arg \max_a Q^\pi(s, a) \ \forall s$?
- What is new ϵ -greedy policy, if $k = 3, \epsilon = 1/k$? Give an example for $\pi(s_1)$.

| s_1 | s_2 | s_3 | s_4 | s_5 | s_6 | s_7 |
|-------|-------|-------|-------|---|-------|-------|
| | | | |  | | |

Today's Agenda


1. Review of Previous Lectures
2. Model-Free Control (Monte Carlo)
- 3. Model-Free Control (Temporal Difference)**
4. Model-Free Control (Q-Learning)

Why TD?

- Temporal-difference (TD) learning has several advantages over Monte-Carlo (MC)
 - Lower variance
 - Online learning
 - Incomplete sequences
- Natural idea: use TD instead of MC in our control loop
 - Apply TD to $Q(s, a)$
 - Use ϵ -greedy policy improvement
 - Update every time-step

Updating Action-Value Functions with SARSA


- Stat, Action, Reward, Next State, Next Action (SARSA) is the TD methods that can be used to evaluate the Q-value



$$Q(S, A) \leftarrow Q(S, A) + \alpha (R + \gamma Q(S', A') - Q(S, A))$$

SARSA For On-Policy Control

-
- 1: Set initial ϵ -greedy policy π , $t = 0$, initial state $s_t = s_0$
 - 2: Take $a_t \sim \pi(s_t)$ // Sample action from policy
 - 3: Observe (r_t, s_{t+1})
 - 4: **loop**
 - 5: Take action $a_{t+1} \sim \pi(s_{t+1})$
 - 6: Observe (r_{t+1}, s_{t+2})
 - 7: $Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(r_t + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t))$
 - 8: $\pi(s_t) = \arg \max_a Q(s_t, a)$ w.prob $1 - \epsilon$, else random
 - 9: $t = t + 1$
 - 10: **end loop**
-



Every **time-step**:

Policy evaluation **Sarsa**, $Q \approx q_\pi$

Policy improvement ϵ -greedy policy improvement

Does SARSA converge? Does it converge for any choice of α ?

Convergence Properties of SARSA

Theorem

SARSA for finite-state and finite-action MDPs converges to the optimal action-value, $Q(s, a) \rightarrow Q^*(s, a)$, under the following conditions:



- 1 The policy sequence $\pi_t(a|s)$ satisfies the condition of GLIE
- 2 The step-sizes α_t satisfy the Robbins-Munro sequence such that

$$\sum_{t=1}^{\infty} \alpha_t = \infty$$

$$\sum_{t=1}^{\infty} \alpha_t^2 < \infty$$


- For ex. $\alpha_t = \frac{1}{t}$ satisfies the above condition.
- Would one want to use a step size choice that satisfies the above in practice? **Likely not.**

Windy Gridworld Example




- Reward = -1 per time-step until reaching goal
- Undiscounted $\gamma = 1$

SARSA on the Windy Gridworld



SARSA Example

```
1: Set initial  $\epsilon$ -greedy policy  $\pi$ ,  $t = 0$ , initial state  $s_t = s_0$ 
2: Take  $a_t \sim \pi(s_t)$  // Sample action from policy
3: Observe  $(r_t, s_{t+1})$ 
4: loop
5:   Take action  $a_{t+1} \sim \pi(s_{t+1})$ 
6:   Observe  $(r_{t+1}, s_{t+2})$ 
7:    $Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(r_t + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t))$ 
8:    $\pi(s_t) = \arg \max_a Q(s_t, a)$  w.prob  $1 - \epsilon$ , else random
9:    $t = t + 1$ 
10: end loop
```

| s_1 | s_2 | s_3 | s_4 | s_5 | s_6 | s_7 |
|-------|-------|-------|-------|---|-------|-------|
| | | | |  | | |

- Initialize $\gamma = 1$, $\epsilon = 1/k$, $k = 1$, and $\alpha = 0.5$
- Initialize $Q(-, a_1) = [1, 0, 0, 0, 0, 0, 10]$, $Q(-, a_2) = [1, 0, 0, 0, 0, 0, 5]$
- Assume we observe tuple by acting in the word according to SARSA:
 $(s_6, a_1, 0, s_7, a_2)$
- What is $Q(s_6, a_1)$ after a SARSA update?

Today's Agenda

1. Review of Previous Lectures
2. Model-Free Control (Monte Carlo)
3. Model-Free Control (Temporal Difference)
4. **Model-Free Control (Q-Learning)**

Off-Policy Learning

- Evaluate *target policy* $\pi(a|s)$ to compute $v^\pi(s)$ or $q^\pi(s, a)$
- While following *behavior policy* $\mu(a|s)$

$$\{s_1, a_1, r_2, \dots, s_T\} \sim \mu$$

- Why is this important?
 - Learn from observing humans or other agents
 - Re-use experience generated from old policies $\pi_1, \pi_2, \dots, \pi_{t-1}$
 - Learn about optimal policy, while following exploratory policy
 - Learn about multiple policies while following one policy

Q-Learning: Learning the Optimal State-Action Value

- SARSA is an on-policy learning algorithm
 - SARSA estimates the value of the current behavior policy (policy used to take actions in the world)
 - i.e., both *target policy* and *behavior policy* are the same
- For MDP:
 - We know that optimal policy π^* (i.e., *target policy*) is deterministic (i.e., greedy)
 - But we need the *behavior policy* to be stochastic (i.e., ϵ -greedy) to explore
- Can we directly estimate the value of the *greedy target policy*, while acting with an ϵ -greedy *behavior policy*?
 - Yes! Q-learning, an off-policy RL algorithm

Q-Learning: Learning the Optimal State-Action Value

Q-learning: off-policy learning of action-values $Q(s, a)$

- Next action is chosen using ϵ -greedy behavior policy $a_{t+1} \sim \mu(\cdot | s_t)$ where

$$\mu(a|s) = \begin{cases} \epsilon/m + 1 - \epsilon, & \text{if } a = \operatorname{argmax}_{a' \in \mathcal{A}} Q(s, a') \\ \epsilon/m, & \text{otherwise} \end{cases}$$

- But when updating Q-values consider alternative successor action according to *greedy target policy* $a'' \sim \pi(\cdot | s_t)$ where

$$\pi(a|s) = \begin{cases} 1, & \text{if } a = \operatorname{argmax}_{a' \in \mathcal{A}} Q(s, a') \\ 0, & \text{otherwise} \end{cases}$$

or since this is deterministic:

$$a'' = \pi(s) = \operatorname{argmax}_{a' \in \mathcal{A}} Q(s, a')$$

Q-Learning: Learning the Optimal State-Action Value

$$a_{t+1} \sim \mu(\cdot | s_t)$$

$$a'' \sim \pi(\cdot | s_t) \text{ or since deterministic } a'' = \pi(s_t) = \arg \max_{a' \in \mathcal{A}} Q(s_t, a')$$

SARSA Update:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(r_t + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t))$$

Q-Learning Update:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(r_t + \gamma Q(s_{t+1}, a'') - Q(s_t, a_t))$$

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(r_t + \gamma Q(s_{t+1}, \arg \max_{a' \in \mathcal{A}} Q(s_{t+1}, a')) - Q(s_t, a_t))$$

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(r_t + \gamma \max_{a' \in \mathcal{A}} Q(s_{t+1}, a') - Q(s_t, a_t))$$

Q-Learning Algorithm

-
- 1: Initialize $Q(s, a), \forall s \in S, a \in A$ $t = 0$, initial state $s_t = s_0$
 - 2: Set π_b to be ϵ -greedy w.r.t. Q
 - 3: **loop**
 - 4: Take $a_t \sim \pi_b(s_t)$ // Sample action from policy
 - 5: Observe (r_t, s_{t+1})
 - 6: $Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(r_t + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t))$
 - 7: $\pi(s_t) = \arg \max_a Q(s_t, a)$ w.prob $1 - \epsilon$, else random
 - 8: $t = t + 1$
 - 9: **end loop**
-


Q-Learning Convergence

- What conditions are sufficient to ensure that Q-learning with ϵ -greedy exploration converges to **optimal Q^*** ?
 - Visit all (s, a) pairs infinitely often, and the step-sizes α_t satisfy the Robbins-Munro sequence. Note: the algorithm does not have to be greedy in the limit of infinite exploration (GLIE) to satisfy this (could keep ϵ large).
- What conditions are sufficient to ensure that Q-learning with ϵ -greedy exploration converges to **optimal π^*** ?
 - The algorithm is GLIE, along with the above requirement to ensure the Q value estimates converge to the optimal Q.

Q-learning Example

-
- 1: Initialize $Q(s, a), \forall s \in S, a \in A$ $t = 0$, initial state $s_t = s_0$
 - 2: Set π_b to be ϵ -greedy w.r.t. Q
 - 3: **loop**
 - 4: Take $a_t \sim \pi_b(s_t)$ // Sample action from policy
 - 5: Observe (r_t, s_{t+1})
 - 6: $Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(r_t + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t))$
 - 7: $\pi(s_t) = \arg \max_a Q(s_t, a)$ w.prob $1 - \epsilon$, else random
 - 8: $t = t + 1$
 - 9: **end loop**
-

- Initialize $\gamma = 1, \epsilon = 1/k, k = 1$, and $\alpha = 0.5$
- Initialize $Q(-, a_1) = [1, 0, 0, 0, 0, 0, 10]$, $Q(-, a_2) = [1, 0, 0, 0, 0, 0, 5]$
- Assume we observe tuple by acting in the word: $(s_6, a_1, 0, s_7)$
- What is $Q(s_6, a_1)$ after a Q-Learning update? How does this compare to SARSA?

| s_1 | s_2 | s_3 | s_4 | s_5 | s_6 | s_7 |
|-------|-------|-------|-------|---|-------|-------|
| | | | |  | | |

Cliff Walking Example

- Q-learning is more optimistic than SARSA by taking the max action
- SARSA is preferred in environments where you must be more cautious
 - Environments with large negative rewards
 - A real robot acting in real world, which might break or cause damage
- Q-Learning is better in environments where you do not need to be cautious
 - Where we do not have many negative results
 - Training in simulation environment
- Both algorithms eventually converge

