



REINFORCEMENT LEARNING

CP8319/CPS824

Lecture 12

Instructor: Nariman Farsad

Today's Agenda

- 1. Review of Previous Lectures**
2. Value Function Approximation Policy Evaluation


Policy Iteration: Known Model

Set $i = 0$

Initialize $\pi_0(s)$ randomly for all states s

While $i \neq 0$ or $\|\pi_i - \pi_{i-1}\|_1 > 0$ (L1-norm, measures if the policy changed for any state):

- $v^{\pi_i} \leftarrow$ MDP value function **policy evaluation** of π_i
- $\pi_{i+1} \leftarrow$ **Policy improvement** on v^{π_i}
- $i = i + 1$



Policy evaluation Estimate v_π

Iterative policy evaluation

Policy improvement Generate $\pi' \geq \pi$

Greedy policy improvement

Model Free Policy Iteration

Set $i = 0$

Initialize $\pi_0(s)$ randomly for all states s

While $i \neq 0$ or $\|\pi_i - \pi_{i-1}\|_1 > 0$ (L1-norm, measures if the policy changed for any state):


- $Q^{\pi_i} \leftarrow$ MDP value function **MC policy Q evaluation** of π_i
- $\pi_{i+1} \leftarrow$ **ϵ -greedy Policy improvement** on Q^{π_i}
- $i = i + 1$

greedy(Q)

$$\pi_{i+1}(a|s) = \begin{cases} 1, & \text{if } a = \operatorname{argmax}_{a' \in \mathcal{A}} Q^{\pi_i}(s, a') \\ 0, & \text{otherwise} \end{cases}$$

ϵ -greedy(Q)

$$\pi_{i+1}(a|s) = \begin{cases} \epsilon/m + 1 - \epsilon, & \text{if } a = \operatorname{argmax}_{a' \in \mathcal{A}} Q^{\pi_i}(s, a') \\ \epsilon/m, & \text{otherwise} \end{cases}$$



Policy evaluation Monte-Carlo policy evaluation, $Q = q_\pi$

Policy improvement ϵ -greedy policy improvement

Greedy in the Limit of Infinite Exploration (GLIE)

Definition of GLIE

- All state-action pairs are visited an infinite number of times

$$\lim_{i \rightarrow \infty} N_i(s, a) \rightarrow \infty$$

- Behavior policy (policy used to act in the world) converges to greedy policy

$$\lim_{i \rightarrow \infty} \pi(a|s) \rightarrow \arg \max_a Q(s, a) \text{ with probability 1}$$


- A simple GLIE strategy is ϵ -greedy where ϵ is reduced to 0 with the following rate: $\epsilon_i = 1/i$

Monte Carlo Online Control/On Policy Improvement

```
1: Initialize  $Q(s, a) = 0, N(s, a) = 0 \forall (s, a)$ , Set  $\epsilon = 1, k = 1$ 
2:  $\pi_k = \epsilon$ -greedy( $Q$ ) // Create initial  $\epsilon$ -greedy policy
3: loop
4:   Sample  $k$ -th episode  $(s_{k,1}, a_{k,1}, r_{k,1}, s_{k,2}, \dots, s_{k,T})$  given  $\pi_k$ 
4:    $G_{k,t} = r_{k,t} + \gamma r_{k,t+1} + \gamma^2 r_{k,t+2} + \dots + \gamma^{T-t-1} r_{k,T}$ 
5:   for  $t = 1, \dots, T$  do
6:     if First visit to  $(s, a)$  in episode  $k$  then
7:        $N(s, a) = N(s, a) + 1$ 
8:        $Q(s_t, a_t) = Q(s_t, a_t) + \frac{1}{N(s,a)} (G_{k,t} - Q(s_t, a_t))$ 
9:     end if
10:  end for
11:   $k = k + 1, \epsilon = 1/k$ 
12:   $\pi_k = \epsilon$ -greedy( $Q$ ) // Policy improvement
13: end loop
```

SARSA For On-Policy Control

-
- 1: Set initial ϵ -greedy policy π , $t = 0$, initial state $s_t = s_0$
 - 2: Take $a_t \sim \pi(s_t)$ // Sample action from policy
 - 3: Observe (r_t, s_{t+1})
 - 4: **loop**
 - 5: Take action $a_{t+1} \sim \pi(s_{t+1})$
 - 6: Observe (r_{t+1}, s_{t+2})
 - 7: $Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(r_t + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t))$
 - 8: $\pi(s_t) = \arg \max_a Q(s_t, a)$ w.prob $1 - \epsilon$, else random
 - 9: $t = t + 1$
 - 10: **end loop**
-



Every **time-step**:

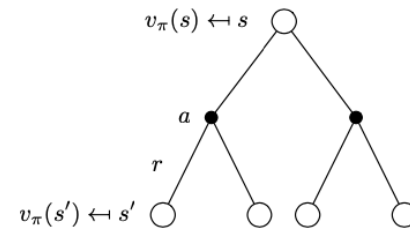
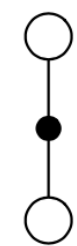
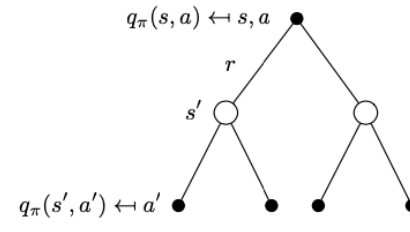
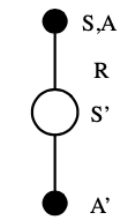
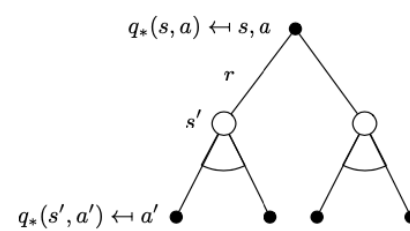
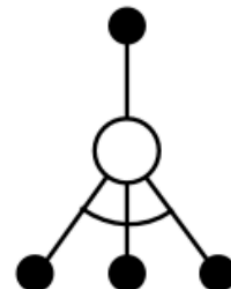
Policy evaluation **Sarsa**, $Q \approx q_\pi$

Policy improvement ϵ -greedy policy improvement

Q-Learning Algorithm

-
- 1: Initialize $Q(s, a), \forall s \in S, a \in A$ $t = 0$, initial state $s_t = s_0$
 - 2: Set π_b to be ϵ -greedy w.r.t. Q
 - 3: **loop**
 - 4: Take $a_t \sim \pi_b(s_t)$ // Sample action from policy
 - 5: Observe (r_t, s_{t+1})
 - 6: $Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(r_t + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t))$
 - 7: $\pi(s_t) = \arg \max_a Q(s_t, a)$ w.prob $1 - \epsilon$, else random
 - 8: $t = t + 1$
 - 9: **end loop**
-

Relationship Between DP (Known) and TD (Unknown)

	<i>Full Backup (DP)</i>	<i>Sample Backup (TD)</i>
Bellman Expectation Equation for $v_{\pi}(s)$	 <p>Iterative Policy Evaluation</p>	 <p>TD Learning</p>
Bellman Expectation Equation for $q_{\pi}(s, a)$	 <p>Q-Policy Iteration</p>	 <p>Sarsa</p>
Bellman Optimality Equation for $q_{*}(s, a)$	 <p>Q-Value Iteration</p>	 <p>Q-Learning</p>

Today's Agenda

1. Review of Previous Lectures
- 2. Value Function Approximation Policy Evaluation**

Large-Scale Reinforcement Learning

- Reinforcement learning can be used to solve large problems, e.g.
 - Backgammon: 10^{20} states
 - Computer Go: 10^{170} states
 - Helicopter: continuous state space
- How can we scale up the model-free methods for prediction and control from the last two lectures?




Value Function Approximation

- So far we have represented value function by a lookup table
 - Every state s has an entry $v(s)$
 - Or every state-action pair s, a has an entry $Q(s, a)$
- Problem with large MDPs:
 - There are too many states and/or actions to store in memory
 - It is too slow to learn the value of each state individually
- Solution for large MDPs:
 - Estimate value function with *function approximation*

$$v^{\pi}(s) \approx \hat{v}(s; \mathbf{w})$$
$$\text{or } Q^{\pi} \approx \hat{Q}(s, a; \mathbf{w})$$

- *Generalize* from seen states to unseen states
- Update parameter \mathbf{w} using MC or TD learning

Types of Value Function Approximation



Function Approximators

- Many possible function approximators including:
 - Linear combinations of features
 - Neural networks
 - Decision trees
 - Nearest neighbors
 - Fourier/ wavelet bases
- In this class we will focus on function approximators that are differentiable (Why?)
- Two very popular classes of differentiable function approximators
 - Linear feature representations (2 lectures)
 - Neural networks (Next 2 lectures)
- We require a training method that is suitable for non-stationary, non-iid data

Gradient Descent

- Let $J(\mathbf{w})$ be a differentiable function of parameter vector \mathbf{w}
- Define the *gradient* of $J(\mathbf{w})$ to be:

$$\nabla_{\mathbf{w}}J(\mathbf{w}) = \begin{pmatrix} \frac{\partial J(\mathbf{w})}{\partial w_1} \\ \vdots \\ \frac{\partial J(\mathbf{w})}{\partial w_n} \end{pmatrix}$$

- To find a *local minimum* of $J(\mathbf{w})$:

$$\mathbf{w} \leftarrow \mathbf{w} - \alpha \nabla_{\mathbf{w}}J(\mathbf{w})$$

Value Function Approximation with an Oracle

- First assume we could query any state s and an oracle would return the true value for $v^\pi(s)$
- Therefore, we could have a set $\{(s_1, v^\pi(s_1)), (s_2, v^\pi(s_2)), \dots\}$ of data
- Goal: Find the parameter vector \mathbf{w} that minimizes the loss between a true value function $v^\pi(s)$ and its approximation $\hat{v}(s; \mathbf{w})$ as represented with a particular function/model class parameterized by \mathbf{w} .
- What does this remind you off?

How do we learn?

Batch Gradient Descent:

- Expensive to compute gradient for large dataset
- Computational and space complexity high

$$\theta = \theta - \eta \cdot \nabla_{\theta} J(\theta)$$

Stochastic Gradient Descent:


- Lots of random motion (slow to converge)

$$\theta = \theta - \eta \cdot \nabla_{\theta} J(\theta; x^{(i)}; y^{(i)})$$

Mini-batch Gradient Descent (Mini-batch of size n):

- Hybrid between the two, still stochastic but less random

$$\theta = \theta - \eta \cdot \nabla_{\theta} J(\theta; x^{(i:i+n)}; y^{(i:i+n)})$$



Value Function Approx. By Stochastic Gradient Descent

- Goal: Find the parameter vector \mathbf{w} that minimizes the loss between a true value function $v^\pi(s)$ and its approximation $\hat{v}(s; \mathbf{w})$ as represented with a particular function/model class parameterized by \mathbf{w} .
- Generally, use mean squared error and define the loss as

$$J(\mathbf{w}) = \mathbb{E}_\pi \left[\left(v^\pi(s) - \hat{v}(s; \mathbf{w}) \right)^2 \right]$$

$$\mathbf{w} \leftarrow \mathbf{w} - \frac{1}{2} \alpha \nabla_{\mathbf{w}} J(\mathbf{w})$$

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha \mathbb{E}_\pi \left[\left(v^\pi(s) - \hat{v}(s; \mathbf{w}) \right) \nabla_{\mathbf{w}} \hat{v}(s; \mathbf{w}) \right]$$

- Stochastic gradient descent (SGD) *samples* the gradient:

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha \left(v^\pi(s) - \hat{v}(s; \mathbf{w}) \right) \nabla_{\mathbf{w}} \hat{v}(s; \mathbf{w})$$

- Expected update is equal to full gradient update

Model Free VFA Policy Evaluation

- Don't actually have access to an oracle to tell true $v^\pi(s)$ for any state s
- Now consider how to do model-free value function approximation for prediction / evaluation / policy evaluation without a model

What we did before

- Recall model-free policy evaluation from prior lectures
 - Following a fixed policy π for sampling
 - Goal is to estimate v^π and/or Q^π
- We did this by maintaining *a look-up table* to store estimates of v^π and/or Q^π
 - Update the look-up table estimate after each episode (Monte Carlo)
 - After each step (temporal difference)
- *Now*: in value function approximation, change the estimate update step to include fitting the function approximator

Feature Vectors

- The state is represented by a *feature vector*

$$\mathbf{x}(s) = \begin{pmatrix} x_1(s) \\ \vdots \\ x_n(s) \end{pmatrix}$$

- For example:
 - Distance of a robot from “landmarks”
 - Trends in the stock market
 - Piece and pawn configurations in chess

Linear Value Function Approximation With Oracle

- Represent a value function (or state-action value function) for a particular policy with a weighted linear combination of features:

$$\hat{v}(s; \mathbf{w}) = \sum_{j=0}^n x_j(s) w_j = \mathbf{x}(s)^T \mathbf{w}$$

- Objective function is:

$$J(\mathbf{w}) = \mathbb{E}_{\pi} \left[\left(v^{\pi}(s) - \hat{v}(s; \mathbf{w}) \right)^2 \right]$$

- Update rule is:

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha \left(v^{\pi}(s) - \hat{v}(s; \mathbf{w}) \right) \nabla_{\mathbf{w}} \hat{v}(s; \mathbf{w})$$

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha \left(v^{\pi}(s) - \hat{v}(s; \mathbf{w}) \right) \mathbf{x}(s)$$

Table Lookup Features

- Table lookup is a special case of linear value function approximation
- Using *table lookup features*

$$\mathbf{x}^{table}(S) = \begin{pmatrix} \mathbf{1}(S = s_1) \\ \vdots \\ \mathbf{1}(S = s_n) \end{pmatrix}$$

- Parameter vector \mathbf{w} gives value of each individual state

$$\hat{v}(S, \mathbf{w}) = \begin{pmatrix} \mathbf{1}(S = s_1) \\ \vdots \\ \mathbf{1}(S = s_n) \end{pmatrix} \cdot \begin{pmatrix} \mathbf{w}_1 \\ \vdots \\ \mathbf{w}_n \end{pmatrix}$$

Monte Carlo Value Function Approximation

- Return G_t is an unbiased but noisy sample of the true expected return $v^\pi(s_t)$
- Therefore, can reduce MC VFA to doing supervised learning on a set of (state,return) pairs: $\langle s_1, G_1 \rangle, \langle s_2, G_2 \rangle, \dots, \langle s_T, G_T \rangle$
 - Substitute G_t for the true $v^\pi(s_t)$ when fit function approximator
- Concretely when using linear VFA for policy evaluation

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha(G_t - \hat{v}(s; \mathbf{w})) \nabla_{\mathbf{w}} \hat{v}(s; \mathbf{w})$$

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha(G_t - \hat{v}(s; \mathbf{w})) \mathbf{x}(s)$$

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha(G_t - \mathbf{x}(s)^T \mathbf{w}) \mathbf{x}(s)$$

- Note: G_t may be a very noisy estimate of true return

MC Linear VFA for Policy Evaluation

```
1: Initialize  $\mathbf{w} = \mathbf{0}$ ,  $k = 1$ 
2: loop
3:   Sample  $k$ -th episode  $(s_{k,1}, a_{k,1}, r_{k,1}, s_{k,2}, \dots, s_{k,L_k})$  given  $\pi$ 
4:   for  $t = 1, \dots, L_k$  do
5:     if First visit to  $(s)$  in episode  $k$  then
6:        $G_t(s) = \sum_{j=t}^{L_k} r_{k,j}$ 
7:       Update weights:

8:     end if
9:   end for
10:   $k = k + 1$ 
11: end loop
```

MC Linear VFA for Policy Evaluation Example

Recall MC update: $\alpha(G_t - \mathbf{x}(s)^T \mathbf{w})\mathbf{x}(s)$


Two actions:

- a_1 goes to 7 (solid line)
- a_2 goes to 1-6 with 1/6 probability (dashed line)

Observe $s_1, a_1, 0, s_7, a_1, 0, s_7, a_1, 0, \text{terminate}$

Assume $\mathbf{w}_0 = [1, 1, 1, 1, 1, 1, 1, 1], \alpha = 0.5, \gamma = 0.9$

What is \mathbf{w}_1 after the first SGD update?



MC Linear VFA for Policy Evaluation Example

Recall MC update: $\alpha(G_t - \mathbf{x}(s)^T \mathbf{w})\mathbf{x}(s)$


Two actions:

- a_1 goes to 7 (solid line)
- a_2 goes to 1-6 with 1/6 probability (dashed line)

Observe $s_1, a_1, 0, s_7, a_1, 0, s_7, a_1, 0, \text{terminate}$

Assume $\mathbf{w}_0 = [1, 1, 1, 1, 1, 1, 1, 1], \alpha = 0.5, \gamma = 0.9$

What is \mathbf{w}_1 after the first SGD update?



Recall: Temporal Difference Learning w/ Lookup Table

- Uses bootstrapping and sampling to approximate $v^\pi(s)$
- Simplest temporal-difference learning algorithm: TD(0)
 - Update value $v^\pi(s_t)$ toward estimated return $r_t + \gamma v^\pi(s_{t+1})$

$$v^\pi(s_t) = v^\pi(s_t) + \alpha([r_t + \gamma v^\pi(s_{t+1})] - v^\pi(s_t))$$

- $r_t + \gamma v^\pi(s_{t+1})$ is called the TD target
- $\delta_t = r_t + \gamma v^\pi(s_{t+1}) - v^\pi(s_t)$ is called the TD error
- a biased estimate of the true value $v^\pi(s)$
- Represent value for each state with a separate table entry

Temporal Difference (TD(0)) Learning with VFA

- In value function approximation, the target is $r + \gamma \hat{v}^\pi(s'; \mathbf{w})$, a biased and approximated estimate of the true value $v^\pi(s)$
- Can reduce doing TD(0) learning with value function approximation to supervised learning on a set of data pairs:

$$\langle s_1, r_1 + \gamma \hat{v}^\pi(s_2; \mathbf{w}) \rangle, \langle s_2, r_2 + \gamma \hat{v}^\pi(s_3; \mathbf{w}) \rangle, \dots$$

- Find weights to minimize mean squared error:

$$J(\mathbf{w}) = \mathbb{E}_\pi \left[\left(r_j + \gamma \hat{v}^\pi(s_{j+1}; \mathbf{w}) - \hat{v}(s_j; \mathbf{w}) \right)^2 \right]$$

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha \left(r_j + \gamma \mathbf{x}(s_{j+1})^T \mathbf{w} - \mathbf{x}(s_j)^T \mathbf{w} \right) \mathbf{x}(s_j)$$

- Therefore, we are using 3 forms of approximation, what are they?

TD(0) Linear VFA for Policy Evaluation

1: Initialize $\mathbf{w} = \mathbf{0}$, $k = 1$

2: **loop**

3: Sample tuple (s_k, a_k, r_k, s_{k+1}) given π

4: Update weights:

$$\mathbf{w} = \mathbf{w} + \alpha(r + \gamma \mathbf{x}(s')^T \mathbf{w} - \mathbf{x}(s)^T \mathbf{w}) \mathbf{x}(s)$$

5: $k = k + 1$

6: **end loop**

TD(0) Linear VFA for Policy Evaluation Example

Recall MC update: $\alpha(G_t - \mathbf{x}(s)^T \mathbf{w})\mathbf{x}(s)$

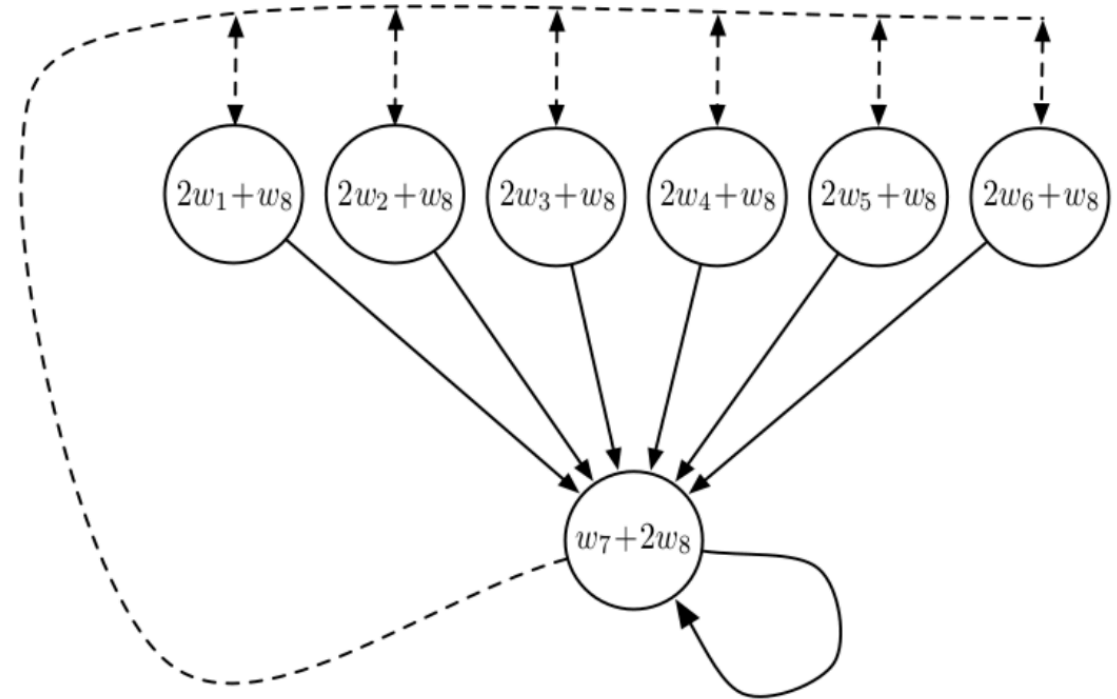
Two actions:

- a_1 goes to 7 (solid line)
- a_2 goes to 1-6 with 1/6 probability (dashed line)

Observe $(s_1, a_1, 0, s_7)$

Assume $\mathbf{w}_0 = [1, 1, 1, 1, 1, 1, 1, 1]$, $\alpha = 0.5$, $\gamma = 0.9$

What is \mathbf{w}_1 after the first SGD update?



TD(0) Linear VFA for Policy Evaluation Example

Recall MC update: $\alpha(G_t - \mathbf{x}(s)^T \mathbf{w})\mathbf{x}(s)$


Two actions:

- a_1 goes to 7 (solid line)
- a_2 goes to 1-6 with 1/6 probability (dashed line)

Observe $(s_1, a_1, 0, s_7)$

Assume $\mathbf{w}_0 = [1, 1, 1, 1, 1, 1, 1, 1]$, $\alpha = 0.5$, $\gamma = 0.9$

What is \mathbf{w}_1 after the first SGD update?



Convergence Rates for Linear Value Function Approximation for Policy Evaluation

- Does TD or MC converge faster to a fixed point?
- Not (to my knowledge) definitively understood
- Practically TD learning often converges faster to its fixed value function approximation point